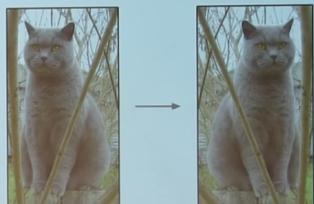


# practical advice for working with images

## Data Augmentation

### 1. Horizontal flips



## Data Augmentation

### 2. Random crops/scales

**Training:** sample random crops / scales

ResNet:

1. Pick random L in range [256, 480]
2. Resize training image, short side = L
3. Sample random 224 x 224 patch



**Testing:** average a fixed set of crops

ResNet:

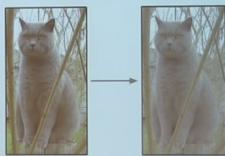
1. Resize image at 5 scales: {224, 256, 384, 480, 640}
2. For each size, use 10 224 x 224 crops: 4 corners + center, + flips

## Data Augmentation

### 3. Color jitter

#### Simple:

Randomly jitter contrast



#### Complex:

1. Apply PCA to all [R, G, B] pixels in training set
2. Sample a "color offset" along principal component directions
3. Add offset to all pixels of a training image

(As seen in [Krizhevsky et al. 2012], ResNet, etc)

## Data Augmentation

### 4. Get creative!

Random mix/combinations of :

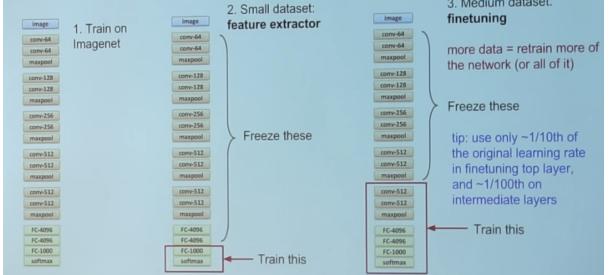
- translation
- rotation
- stretching
- shearing,
- lens distortions, ... (go crazy)

## Transfer Learning

"You need a lot of data if you want to train your CNNs"

**BUSTED**

## Transfer Learning with CNNs

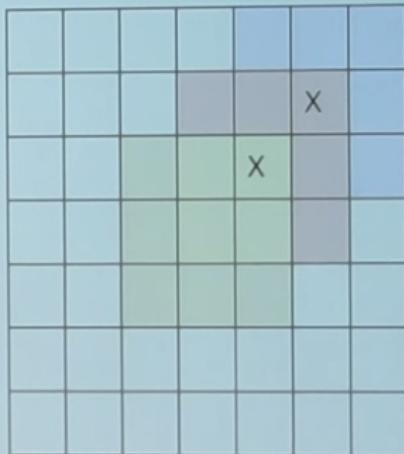


more generic	very similar dataset	very different dataset
more specific		
<p>The diagram shows a vertical stack of layers. From top to bottom: 'image' (blue), 'conv-64', 'conv-64', 'maxpool', 'conv-128', 'conv-128', 'maxpool', 'conv-256', 'conv-256', 'maxpool', 'conv-512', 'conv-512', 'maxpool', 'conv-512', 'conv-512', 'maxpool', 'FC-4096', 'FC-4096', 'FC-1000', 'softmax' (green). Arrows point from the first two layers to the first column of the table, and from the last three layers to the second column.</p>	<b>very little data</b> Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
<b>quite a lot of data</b>	Finetune a few layers	Finetune a larger number of layers

## The power of small filters

**Question:** If we stack **three** 3x3 conv layers, how big of an input region does a neuron in the third layer see?

**Answer:** 7 x 7



Three 3 x 3 conv gives similar representational power as a single 7 x 7 convolution

## The power of small filters

Suppose input is  $H \times W \times C$  and we use convolutions with  $C$  filters to preserve depth (stride 1, padding to preserve  $H, W$ )

one CONV with  $7 \times 7$  filters

Number of weights:  
 $= C \times (7 \times 7 \times C) = 49 C^2$

Number of multiply-adds:  
 $= (H \times W \times C) \times (7 \times 7 \times C)$   
 $= 49 HWC^2$

three CONV with  $3 \times 3$  filters

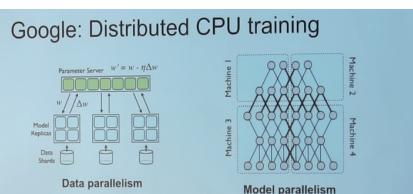
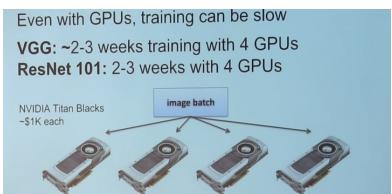
Number of weights:  
 $= 3 \times C \times (3 \times 3 \times C) = 27 C^2$

Number of multiply-adds:  
 $= 3 \times (H \times W \times C) \times (3 \times 3 \times C)$   
 $= 27 HWC^2$

- Also, if we stick ReLUs in between these more layers but smaller convs, we get more non-linearity in our model for the same computations.
- This is the theme of tricks used in inception modules. Save on number of params and computations while getting more expressive power!
- We can take convolutional layers and express them as matrix multiplies. this is really good from the perspective of taking advantage of matrix multiply optimizations that have been built for various computing platforms, GPUs being the main ones.

## GPUs can be programmed

- CUDA (NVIDIA only)
  - Write C code that runs directly on the GPU
  - Higher-level APIs: cuBLAS, cuFFT, cuDNN, etc
- OpenCL
  - Similar to CUDA, but runs on anything
  - Usually slower :(



GPU - CPU communication is a bottleneck.  
=>

CPU data prefetch+augment thread running  
while  
GPU performs forward/backward pass

### Floating point precision

- 64 bit "double" precision is default in a lot of programming
- 32 bit "single" precision is typically used for CNNs for performance

## Implementation details: Recap

- GPUs much faster than CPUs
- Distributed training is sometimes used
  - Not needed for small problems
- Be aware of bottlenecks: CPU / GPU, CPU / disk
- Low precision makes things faster and still works
  - 32 bit is standard now, 16 bit soon
  - In the future: binary nets?